



P34

The v7.6.0 Evidence Ledger and Prediction Update: Three-State Audit, Downgraded Claims, and F6–F10

v7.6.0 文献证据台账与预测更新

TwinsEarth · 智能体宇宙论文系列



The v7.6.0 Evidence Ledger and Prediction Update: Three-State Audit, Downgraded Claims, and F6–F10

Working Title (EN): The v7.6.0 Evidence Ledger and Prediction Update: Three-State Literature Audit, Downgraded Claims, and Falsifiable Forecasts F6–F10

Whole-volume closing paper P34 • upgrading the pre-verification shard's EVIDENCE_LEDGER_v76.md into the formal paper-version evidence ledger • updating the hard predictions of P15/P33 • UDOS Reasoning Engine v7.6.0 • Fang Wenxin • 2026-09-21

One sentence first: not one external reference cited in the whole v7.6.0 volume is "heard about"—each was verified online, and pinned to one of three states: verified verified, partially verified partially verified (giving actual values), no traceable evidence no traceable evidence (marking [CITATION NEEDED]). This paper makes this ledger public, and makes public **which widely circulating numbers must be downgraded or**

removed (MAPF million success rate, five attributes of settlement neutrality, Hybrid 96.1%, etc.), finally retaining P15's F1–F5 and adding five hard predictions F6–F10 scorable in 2027.

Evidence caliber (stated only once in the whole paper): this paper is an **evidence ledger and prediction registration** paper, containing no new experiments, no new UDOS measured numbers. The three-state verification was completed **2026-09-21** (Asia/Shanghai), the method being independently reviewing each clue with arXiv abs/html, ACL Anthology, IEEE Xplore, OpenReview, Zenodo, official press releases; the arXiv numbers, figures, institutional attributions the user pasted are all treated as "clues" rather than facts. A verified mark only represents "the reference truly exists, title/year/source consistent," its in-paper numbers still marked "paper report • not independently reviewed"; a partially verified mark is subject to the **actual values** this paper gives; a no traceable evidence mark is always written [CITATION NEEDED], forbidden to write as fact.

Abstract

The v7.6.0 paper volume (P16–P34) cites about 50 external clues. This paper does not do survey stacking, but runs all clues through **three-state verification**: verified verified (exists and bibliographic consistent), partially verified partially verified (exists but some number/record number inconsistent with the clue, giving actual values), no traceable evidence no traceable evidence (marking [CITATION NEEDED]). The verification conclusion is optimistic but restrained: the core new papers (H-CSC, EPOB, ClawArena-Team, Warp-Cortex, RSI-Agent, Ouroboros, OASIS) **the vast majority truly exist and core propositions hold**; what really needs downgrading is four classes—(a) MAPF's "million-scale 99.9%/163μs" only seen in Russian media relaying; (b) "five attributes of settlement neutrality" not located to any original source; (c) Hybrid LangGraph-CrewAI's 96.1%/–76.2% token not verified to original text; (d) ISONOMIA/RSIAIO/Adaptive Bitmask's record numbers and simulation numbers need correction. This paper also registers a "removed/downgraded list," and on the basis of retaining P15 F1–F5 adds five falsifiable predictions F6–F10 (three-level finality, scope, BRS+DRS, review gate, rule predicate vs learned predicate), all marked "pending 2027 scoring."

Keywords: evidence ledger; three-state verification; citation audit; downgrade list; falsifiable prediction; literature credibility

Structured Abstract (background problem → method/argument → evidence → contribution)

- **Background problem:** the easiest place an AI paper volume fails is not argument but **citation**—a widely circulating arXiv number or "million-scale" figure, if written into text without verification, builds the whole argument on sand.
- **Core argument:** an honest paper volume must pin each external material it cites to one of three states, and make public "which I believed, which I discounted, which I did not find."
- **Evidence clues:** about 50 clues verified online one by one; core new papers verified in majority, partially verified concentrated on record numbers and media-relayed figures, no traceable evidence concentrated on sourceless "framework five attributes" classes.
- **Contribution:** a formal paper-version three-state verification ledger; a removed/downgraded list; P15 F1–F5 retained + F6–F10 added prediction update; and the whole v7.6.0 volume's 19-paper index.

1 Why a Public Ledger Is Needed

AI writing's default failure mode is: the model taking "impressions" seen in training as "facts," taking second-hand relayed numbers as paper original text, taking an unfindable framework as "existing consensus." Since v7.6.0 calls itself "measurement and governance," it must execute the same governance on its own **citations**. Therefore this paper upgrades the pre-verification shard's working ledger into a formal paper, making public the three-state results and downgrade list—this is both responsible to readers and the extension at the literature layer of P12's evidence grading methodology.

The three-state definitions stated only once in the whole paper:

- **Verified verified:** the reference truly exists, title/author/year/source consistent; in-paper numbers still marked "paper report • not independently reviewed."
- **Partially verified partially verified:** the reference exists, but some number/record number/title inconsistent with the clue, this paper writes the **actual value**.
- **No traceable evidence no traceable evidence:** cannot find or seriously inconsistent with the given number, marking [CITATION NEEDED] when writing.




Figure 2 Distribution of v7.6.0 whole-volume external clues classified by three states (conceptual illustration; counts subject to this paper's ledger). Conceptual illustration, not measured data.

2 Core New Paper Ledger (Group A, v7.6 main dialogue objects)

Table 1 Group A: the eight core papers directly in dialogue with v7.6.0

No.	Reference	Status	Actual information verified	Safe citation caliber
A1	H-CSC (semantic commit BFT)	verified (v2 retitled)	v2 retitled Certifiable Semantic Agreement Among LLM Agents: What the Admissibility Instrument Decides (arXiv:2606.07316v2); typed commit objects and 2f+1 certificate envelopes are true. Negative results verbatim hit: aggregation distance 5.32°/5.48° vs steelman 4.27°/4.32°; honestly retaining the author's explicit "we withdraw it"; lexical predicate AUROC 0.865–0.982 beats learned encoders 0.621–0.744.	Use v2 caliber: withdrawal + lexical predicate beats learned; numbers paper report • not independently reviewed.
A2	EPOB (project orchestration benchmark)	verified	formally recorded as OpenReview NtPIgzPt0U (2026-05) + epob.us leaderboard; five dimensions Plan/Assignment/Coordination/Deliverable/Efficiency are true.	Five dimension names true; failure taxonomy paper report caliber.
A3	ClawArena-Team	verified	Benchmarking Subagent Orchestration... (arXiv: 2606.31174); 41 scenarios/258 rounds/pure execution-style scoring without LLM judging is true.	Direction true; "no LLM exceeds 50% permission precision," "hundredfold cost/fourfold total score," "SMS three-term product" not verbatim hit, write "paper report • not independently reviewed," not writing fixed percentages.
A4	Warp-Cortex	verified	Singleton weight sharing + topological synapse, memory complexity $O(N \cdot L) \rightarrow O(1) + O(N \cdot k)$ is true.	"100 Agents/2.2GB" paper report; "million" paper self-described theoretical extrapolation, actual measured only about hundred scale.
A5	RSIAgent	verified	curricula/actor/verifier three roles, broad-then-deep strategy true (arXiv:2609.15364).	"open-source Kimi-K3/GLM-5.3 beats GPT-6-class closed" paper/media report; title "New Environments" means learning in new environments, not cross-environment transfer.
A6	Ouroboros	verified	reviewed commits becoming runtime is true (arXiv:2608.08311, v3).	Terminal-Bench 2.1 original 86.97%/after removing reward-hack 86.74% self-reported SOTA, not independently reviewed.
A7	OASIS	verified	million-user social simulation, reproducing information diffusion/polarization/herding is true (arXiv:2411.11581, NeurIPS 2024).	Evidence nature is phenomenon reproduction, not causal identification.
A8	MAPF-GPT-DDG	partially verified	true title Advancing Learnable MAPF Solvers with Active Fine-Tuning (arXiv:2506.23793, IROS 2025), DDG active fine-tuning is true.	"2048×2048 / 524k 100% / 1M 99.9% / 163μs" only seen in Russian media relaying, not accepted as paper fact;

No.	Reference	Status	Actual information verified	Safe citation caliber
				predecessor only small maps, 16–32 Agents.

3 RSI Lineage Ledger (Group B) and Market Governance Ledger (Group C)

Table 2 Group B: RSI systems used for P24 level positioning

No.	Reference	Status	Actual information
B1	The Last AI Built by Humans	verified	arXiv:2609.11873; HCI + seven RSI discrimination criteria, as P24 yardstick.
B2	Dream-RSI	verified	accurate title Recursive Self-Improvement through Evolving Worlds (arXiv: 2609.14858); "317 calls $\approx$ 51,200 generations" paper/official report.
B3	Gödel Agent	verified	ACL 2025, 2025.acl-long.1354 (arXiv:2410.04444); self-referential code-modification self-improvement, second-hand material marks L4, not calling L5.
B4	HyperAgents	verified	arXiv:2603.19461 (Meta FAIR); the meta-level modification process itself editable, closest to "recursive meta-improvement" but not reaching L5.
B5	RSIAIO	partially verified	correctly recorded as Zenodo 15644670 (not user-given 15646184); VM controlled execution + Darwinian Mode.

Table 3 Group C: Markets and governance

No.	Reference	Status	Actual information
C1	AgentLance	verified	arXiv:2608.23867; private-cost bidding + self-maintained strategy notes + VCG-style payment.
C2	Agent Exchange	verified	arXiv:2507.03904; multi-attribute auctions + generalized second price, greedy/random degradation when liquidity insufficient.
C3	ISONOMIA	partially verified	correctly recorded Zenodo 21338480 (not 21287289); mutual-trust bookkeeping unit ergs + bicameral governance true; "45,000 simulations/300 parameter points" not verified, mark [RESULT NEEDED] .
C4	When Agent Markets Arrive	verified	arXiv:2604.06688; first-price sealed-bid auction.
C5	Five attributes of settlement neutrality	no traceable evidence	not located to any original paper/Zenodo record; forbidden to cite as existing framework, rewrite as UDOS P11/P25-authored design goal or mark [CITATION NEEDED] .

4 Communication Scale, Evaluation Frameworks, and Classic/Industry Ledgers (Groups D/E/F/G summary)

Table 4 Group D Communication and scale

No.	Reference	Status	Actual information
D1	RocketMQ-A2A	verified	ACM FSE 2026 Industry (Alibaba Cloud); session-level replayable event streams, LiteTopic session isolation, resume from breakpoint.
D2	Adaptive Bitmask	partially verified	npm product real, sub-10ms direction credible; "85× compression/8.2ms p99/N=5000" not independently reviewed, mark author claim.
D3	AI Agent Protocols survey	verified	arXiv:2504.16736; CNS=(N_success/N_attempts)/T_avg verbatim hit.

Table 5 Group E Evaluation and frameworks

No.	Reference	Status	Actual information
E1	TUMIX	verified	arXiv:2510.01279; highest +3.55% on Gemini-2.5-Pro verbatim hit.
E2	MAKER	verified	arXiv:2511.09030; extreme decomposition + k=3 majority voting, first zero-error completion of over-million-step tasks.
E3	Thinking vs Doing	verified	arXiv:2506.07976, NeurIPS 2025 WS; Gemma 3 12B reaches open-source SOTA on WebVoyager/WebArena.
E4	Hybrid LangGraph-CrewAI	partially verified	hybrid direction truly exists, but IEEE Access 11481053, "96.1%/−76.2% token/14.5× latency" not located to original text, mark pending.
E5	MultiAgentBench	verified	ACL 2025, 2025.acl-long.421; systematically comparing star/chain/tree/graph topologies; precise "graph best" ranking not verbatim hit, write "topology strongly task-related."
E6	AgentSociety	verified	arXiv:2502.08691; 10,000+ Agents, about 5 million interactions verbatim hit.
E7	MegaAgent	verified	ACL 2025 Findings, 2025.findings-acl.259; no predefined SOP, national-level policy simulation expanding to 590 Agents.

Group F Classic background (F1–F10, all verified numbers correct): Kaplan 2020 (2001.08361), Chinchilla 2022 (2203.15556), Attention 2017 (1706.03762), AutoGen 2023 (2308.08155), Practical BFT 1999, HotStuff 2019 (1803.05069), Switch Transformers 2021 (2101.03961), Krogh & Vedelsby 1995, Emergence Mirage 2023 (2304.15004), Conformal Prediction 2021 (2107.07511).

Group G Industry rumors (key items verified one by one, all verified, individual numbers partially verified): Kimi K2.6 Agent Swarm (officially says up to 300 sub-Agents parallel), Kimi K3 (arXiv: 2607.24653, 2.8T parameter open-source MoE), Ant Agentar 2.0 (WAIC 2026, nearly 200 job templates), Inspur single rack 384 OCM CPUs supporting 40,000+ Agents, Huawei openJiuwen (dual-dimension RSI, arXiv:2608.27969), Huawei Tao(τ) Law (ISCAS 2026-05-25, six years mass production 381 chips, 2031 equivalent 1.4nm), Anthropic R&D automation (about 26% R&D AI-led, about 30,000 concurrent Agents; "1 billion decisions" partially verified not found, discarded), OpenAI automated research intern → 2028-03 researcher (about 3.1 Agent workdays/person-day), Pillar 3 Fail-Safe & Recovery (Swarm legal abort/drift rollback/cascade containment).

5 The Removed / Downgraded List (writers must comply)

The following nine are circulating claims this verification **actively abandoned or downgraded**. They are not "outrageously wrong," but "writing into text misleads":

1. **C5 Five attributes of settlement neutrality (no traceable evidence)**: no original source. Can only rewrite as UDOS P11/P25-authored design goal, or mark [CITATION NEEDED].
2. **A8 MAPF million numbers (partially verified downgrade)**: DDG active fine-tuning true; "524k 100% / 1M 99.9% / 163μs" only Russian media relaying, P21 can only say direction, million numbers mark media claim.
3. **B5 RSIAI0 record number (partially verified)**: correct is Zenodo **15644670**, not 15646184.
4. **C3 ISONOMIA record number and simulation numbers (partially verified)**: correct is Zenodo **21338480**; "45,000 simulations/300-300 pass" not verified, mark [RESULT NEEDED].
5. **D2 Adaptive Bitmask (partially verified)**: sub-10ms direction credible, "85× compression/8.2ms p99" mark author claim.
6. **E4 Hybrid LangGraph-CrewAI (partially verified)**: "96.1%/−76.2%/14.5×" not located to original text, downgraded to pending.
7. **A3 ClawArena precise thresholds (partially verified)**: direction true, "50% permission precision/hundredfold cost/SMS formula" not verbatim hit, mark paper report • not independently reviewed.
8. **G6 Anthropic "1 billion decisions" (partially verified)**: 26%/30,000 concurrent verified; "1 billion" not found, discarded.
9. **A1 H-CSC version (partially verified)**: must use v2 caliber—has withdrawn "honest retention advantage," lexical predicate beats learned encoders; must not continue v1 positive conclusions.

6 Prediction Update: Retain F1–F5, Add F6–F10

P15's proposed F1–F5 (structural change share, changing structure changes curve, architecture search boundary, calibration-class gain, cross-substrate master curve collapse) are **all retained**, continuing as 2026–2031 scoring objects. This section adds F6–F10, specifically scoring the success/failure of v7.6.0's eight design proposals—they are all written as **falsifiable conditions**, not established conclusions.



Figure 3 The proposal time (2026-09) and expected scoring window (2027) of the five design predictions F6–F10. Conceptual illustration, not measured data.

- **F6 (three-level finality vs binary stopping):** under the same fault-injection suite, three-level typed finality relative to v7.5.0 binary STOP/CONTINUE should have significantly higher completion rate, false-stop rate should not rise; if no measurable difference, then ① roll back binary. **Criterion:** ≥ 30 seeds, 95% CI not crossing zero. Pending 2027 scoring.
- **F7 (scope tightening boundary-crossing rejection):** after TransferBundle adds scope, structural boundary-crossing call rejection rate should significantly rise, while normal work-order availability loss below preset threshold; if availability loss exceeds boundary-crossing benefit, then ② change to on-demand authorization. Pending scoring.
- **F8 (BRS+DRS coverage of blind zones):** the two-stage flywheel relative to single-stage coverage-aware should measurably raise coverage of the identifiability blind zones P6 identified (acceleration scalar class); if no improvement then ③ fall back single stage. Pending scoring.
- **F9 (review gate intercepting silent failures):** the [structural] gate's interception rate of "version-number accident class" silent failures should be significantly higher than the no-gate baseline, while the cost to merge speed acceptable; if idle then ④ downgrade to a hint. Pending scoring.
- **F10 (rule predicate vs learned predicate AUROC difference):** under the same adversarial attack family, deterministic rule predicates' safety discrimination AUROC should be no lower than learned predicates (continuing H-CSC v2 direction); if learned predicates stably and substantially dominate, then ⑤ "rules as base" needs weakening to attack-family-based splitting. Pending scoring.

7 The v7.6.0 Whole Volume 19-Paper Index

Table 6 P16–P34 quick reference

Paper	Title	One-sentence theme	External dialogue object
P16	Typed finality	binary stopping → three-level semantic commit	H-CSC v2
P17	Least-privilege orchestration	TransferBundle adds scope	ClawArena-Team
P18	Broad-deep two-stage flywheel	BRS+DRS submodular collection	RSIAgent
P19	Structural change review gate	MatrixConfig submission must accompany assertions	Ouroboros
P20	Heterogeneity dividend and homogeneity saturation	independent source count $\neq$ head count	OASIS/TUMIX/P9
P21	Million-level coordination end-to-end evidence	MAPF success cannot extrapolate semantic tasks	MAPF-GPT-DDG
P22	Memory wall is not coordination wall	capacity $\neq$ quality $\neq$ coordination	Warp-Cortex
P23	Project-level five-dimensional evaluation	work order → project failure taxonomy	EPOB
P24	RSI lineage measured positioning	six systems land L1–L5, none L5	Theseus/Dream-RSI/Gödel/HyperAgents
P25	Agent task markets	VCG private information and settlement	AgentLance/Exchange/ISONOMIA
P26	Weighted BFT and reputation	weight direction of equal-weight $2f+1$	WBFT/COALITION-VAST/BlockAgents
P27	Large-scale social simulation causal validity	phenomenon reproduction $\neq$ mechanism verification	OASIS/AgentSociety/MegaAgent
P28	Million-level communication protocol	replayable event streams/sub-10ms	RocketMQ-A2A/Bitmask
P29	Test-time scaled collaboration	tool integration and extreme decomposition voting	TUMIX/MAKER/TTI
P30	Hybrid orchestration and framework-task matching	topology selection decision tree calibration	MultiAgentBench/AutoGen
P31	Fail-safe and recovery	legal abort/drift rollback/cascade containment	Pillar 3
P32	Rules before learning	adversarial scenario predicate robustness boundary	H-CSC v2 lesson generalized
P33	v7.6.0 master architecture	eight structural changes unified spec	whole volume convergence
P34	Evidence ledger and prediction update	three-state verification + F6–F10	this paper

8 Drawbacks, Failure Boundaries, and Falsifiable Conditions

8.1 This Ledger's Own Boundaries

- **Three-state verification is "existence verification," not "correctness verification."** A verified mark only represents "this paper truly exists, bibliographic matches," not that its conclusions are right; in-paper numbers always "paper report • not independently reviewed."
- **Verification completed at abstract level.** A3's SMS formula, E4's 96.1% etc. "not verbatim hit" only means this verification did not read at abstract level, not that the paper necessarily lacks them—they should be opened paper by paper to full text for review before submission, rather than killed with one blow.
- **Industry rumors have short timeliness.** Group G (Kimi/Ant/Inspur/Huawei/OpenAI/Anthropic) comes from official pages and mainstream media, may still change after 2026-09-21, cite with dates.

8.2 Falsifiable Conditions

- **C-ledger-1:** if readers open any verified reference's original text and find title/number/core mechanism inconsistent with this paper's ledger, then that item should immediately be re-marked partially verified/no traceable evidence and the safe citation sentence corrected.
- **C-ledger-2:** if F6–F10 are falsified in the 2027 scoring window, then the corresponding design proposals (P16–P19, P32) should roll back per P33's C1–C5, rather than persisting by changing wording.

9 Conclusion

A paper volume's credibility ultimately depends not on how many papers it cites, but on whether it **dares to publicly admit what was not verified, what was discounted**. v7.6.0's honest position is: the core new papers the vast majority true, core propositions hold; but all widely circulating yet unsound numbers (MAPF million success rate, five attributes of settlement neutrality, Hybrid 96.1%, Anthropic 1 billion decisions), this paper all downgrades or removes, and writes the reasons. F6–F10 turn the eight design proposals from "self-feeling good" into "falsifiable propositions waiting to be slapped."

The ledger is not decoration, it is this volume's disclaimer and scoreboard—what is written in must be answered for, what cannot be verified is never taken as fact.

References

UDOS paper volumes (evidence grades marked in text)

1. UDOS v7.5.0: P9 "Quality saturation law," P12 "Reproducible systems engineering driven by evidence grading," P14 "Convergence of RSI," P15 "Life originates from structural organization" (F1–F5).
2. UDOS v7.6.0: P16–P33 all dedicated papers (papers see Table 6).

External references (existence verified per EVIDENCE_LEDGER_v76.md 2026-09-21; numbers paper reports/company disclosures, not independently reviewed)

1. H-CSC v2 (arXiv:2606.07316); EPOB (OpenReview NtPIgzPtOU, 2026); ClawArena-Team (arXiv: 2606.31174); Warp-Cortex (arXiv:2601.01298); RSI-Agent (arXiv:2609.15364); Ouroboros (arXiv:

- 2608.08311); OASIS (arXiv:2411.11581, NeurIPS 2024); MAPF-GPT-DDG (arXiv:2506.23793, IROS 2025).
2. Theseus Labs (arXiv:2609.11873); Dream-RSI (arXiv:2609.14858); Gödel Agent (ACL 2025, long.1354); HyperAgents (arXiv:2603.19461); RSIAI0 (Zenodo 15644670).
 3. AgentLance (arXiv:2608.23867); Agent Exchange (arXiv:2507.03904); ISONOMIA (Zenodo 21338480); When Agent Markets Arrive (arXiv:2604.06688).
 4. RocketMQ-A2A (ACM FSE 2026 Industry); A Survey of AI Agent Protocols (arXiv:2504.16736).
 5. TUMIX (arXiv:2510.01279); MAKER (arXiv:2511.09030); Thinking vs Doing (arXiv:2506.07976); MultiAgentBench (ACL 2025, long.421); AgentSociety (arXiv:2502.08691); MegaAgent (ACL 2025 Findings, 259).
 6. Classic background: Kaplan 2020 (2001.08361); Chinchilla 2022 (2203.15556); Attention 2017 (1706.03762); AutoGen 2023 (2308.08155); Castro & Liskov 1999; HotStuff (1803.05069); Switch Transformers (2101.03961); Krogh & Vedelsby 1995; Emergence Mirage (2304.15004); Conformal Prediction (2107.07511).
 7. Industry disclosures (Group G): Kimi K2.6/K3 (arXiv:2607.24653); Ant Agentar 2.0 (WAIC 2026); Inspur whole rack (2026-07); Huawei openJiuwen (arXiv:2608.27969) and Tao(τ) Law (ISCAS 2026-05-25); Anthropic R&D automation data (2026-09); OpenAI automated research intern (2026-09-06); Cyber Strategy Institute AI SAFE² Pillar 3 (2026-05).
-

Evidence Discipline and Reproduction Notes

- This paper is an **evidence ledger and prediction registration** paper, no new experiments; three-state verification completed 2026-09-21, the method being independently reviewing each clue online.
- Verified only represents reference existence and bibliographic consistency, in-paper numbers still "not independently reviewed"; partially verified always subject to this paper's actual values; no traceable evidence always [CITATION NEEDED].
- The removed/downgraded nine (Section 5) are this volume's hard discipline, when dedicated papers cite related numbers they must not fall back to the user's original clue values.
- F6–F10 are falsifiable scoring conditions for design proposals, marked "pending 2027 scoring"; before submission partially verified items should be opened paper by paper to full text for secondary review.