



P36

**The 2026 Frontier Influence
Ranking of Multi-Agent Systems
and Recursive
Self-Improvement: An
Evidence-Verified Tally of 20
Papers and 10 Industry
Milestones**



The 2026 Frontier Influence Ranking of Multi-Agent Systems and Recursive Self-Improvement: An Evidence-Verified Tally of 20 Papers and 10 Industry Milestones

Working Title (EN): The 2026 Frontier Influence Ranking of Multi-Agent Systems and Recursive Self-Improvement: An Evidence-Verified Tally of 20 Papers and 10 Industry Milestones

Supplement volume P36 • upgrading EVIDENCE_LEDGER_v771.md 's 20 papers and 10 industry milestones into a dedicated **influence ranking** paper • UDOS Reasoning Engine v7.7.1 • Fang Wenxin • 2026-09-21

Evidence caliber (stated once for the whole paper): this paper is an **editorial influence ranking + evidence verification** paper, containing no new experiments, no new UDOS measured numbers, all external facts taking EVIDENCE_LEDGER_v771.md as the sole fact source. Verification base date **2026-09-21 (Asia/Shanghai, UTC+8)**, method for each clue using arXiv abs/html, ACL Anthology, Zenodo, TechRxiv,

vendor official blogs/sites, industry media first-hand recheck; the list's titles, numbers, figures, institutional attributions all treated as "clues," not facts. Three-state legend:  [existence verified]—searchable, bibliographic record consistent with source or corrected;  [paper/report • not independently rechecked]—existence confirmed, but specific numbers (accuracy/VRAM/concurrency/percentages/share) paper self-reported or company/third-party disclosed, unable to independently recompute;  [no evidence found/unable to independently confirm]—unsearchable or numbers mismatched, **must not be written as fact**. All  numbers uniformly carry "self-reported / company disclosed / media retold" qualifiers.

0 One-Sentence Conclusion + Structured Abstract

One sentence first: in the 2026 multi-Agent and RSI frontier, what truly "lays foundations" is not the highest benchmark, but **the few papers rewriting the paradigm clearly** (MegaAgent's SOP-free autonomy, Warp-Cortex's consumer-grade million-Agent vision, The Last AI Built by Humans' RSI roadmap); but in this ranking **two are empty titles with no evidence** (P1, P17), and another batch of widely circulated numbers (85.5% consensus collapse, 98% context compression, 150+ supporting organizations, 17.8%/19.3% share) can only be used discounted as "self-reported/retold"—ranking is editorial judgment, not citation metrics, evidence strength another axis that must be overlaid.

Structured Abstract (background problem → method/argument → evidence → contribution)

- **Background problem:** multi-Agent and RSI circles produce large numbers of papers and vendor milestones yearly, "who is important" "who is real" mixed together; if ranked only by citations or heat, single-author thin preprints, news columns, company self-statements and formal peer-reviewed papers land on the same tier.
- **Method/argument:** this paper gives each material two independent scores—**influence tier** (foundational ★★★★★ / key driver ★★★★☆ / important supplement ★★★☆☆) and **evidence three-state** (  ). Influence is **editorial judgment**, evidence strength is **verifiable fact**, the two axes stated separately, not contaminating each other.
- **Evidence lines:** 20 papers + 10 industry milestones item by item through the ledger;  about 60%,  concentrated in "exists but numbers self-reported/venue discrepancies,"  concentrated in two sourceless empty titles.
- **Contribution:** one layered 20-paper ranking table, one 10-industry-milestone ranking table; MegaAgent "one paper not two" conflict resolution; no-evidence section; and a group of popular narratives **falsified by the frontier literature itself** (debate always more accurate, Scaling seamlessly extrapolates, RSI already happening).

1 Ranking Method: Two Axes Separated, Editorial Judgment ≠ Citation Metrics

This paper emphasizes three things:

1. **Influence tier is editorial judgment, not citation metrics.** Three tiers set by "how many subsequent works' default assumptions it rewrote," not by citations, downloads, media volume.
★★★★★=foundational (rewriting paradigm or proposing repeatedly citable coordinates);
★★★★☆=key driver (providing key evidence/tools/protocols on some main line);

★★★☆☆=important supplement (providing valuable qualifications or negative results on细分 problems).

2. **Evidence three-state overlaid on tiers, not changing tiers but limiting how to use.** A paper can be ★★★★★ foundational while 🟡 (numbers self-reported)—meaning its **paradigm claims citable**, but its **specific benchmarks only citable as self-reported**; a 🟡 paper however grandly described in clues, **must not be cited as fact**.
3. **Numbers uniformly carry source qualifiers.** "Self-reported"=paper authors measured themselves, no third-party reproduction; "company disclosed"=vendor official caliber; "media retold"=third-party reports passed through media, original not consulted. The three's evidence strength decreases in order.

Section summary: do not read "ranked what" as "trust how much." Tier answers "is it important," three-state answers "can its numbers be used directly"—two sets of questions, this paper answers separately.

2 The 20-Paper Influence Ranking Table

The table below grouped by three tiers, within groups by P number. The **verification status** column directly determines how each item is cited; **core claim** only writes ledger-confirmed or corrected content, self-reported numbers marked in place.

Table 1 20 papers • influence three tiers × evidence three states

Rank	Title (corrected)	Source and number	Core claim (with evidence qualification)	Verification status	UDOS dialogue paper
Tier one ★★★★★ Foundational					
P1	Toward Reliable Collective Intelligence in LLM-MAS (survey: role differentiation/ communication/ shared memory)	Clue claims IEEE Xplore 2026-05; two search rounds no such record	This survey "exists on IEEE Xplore 2026-05"— no evidence found , only thematically adjacent surveys, no same-named literature	■	P36 (must disclose honestly)
P2	Warp-Cortex: An Asynchronous, Memory-Efficient Architecture for Million-Agent Cognitive Scaling on Consumer Hardware	arXiv:2601.01298v1 (2026-01-03, single author Jorge L. Ruiz Williams, no institutional attribution)	Singleton Weight Sharing + Topological Synapse, O(1) weights / O(N • k) context, single RTX 4090 100 concurrent @2.2GB VRAM, theoretical >1000, Referential Injection— all single-author self-reported ; "98% context compression no semantic loss" not in abstract	●	P0 / P14
P3	MegaAgent: A Large-Scale Autonomous LLM-based Multi-Agent System Without Predefined SOPs (early alias ...A Practical Framework...)	ACL 2025 Findings, 2025.findings-acl. 259, pp.4998–5036; arXiv:2408.09955	Large-scale autonomous MAS without predefined SOPs, scaled to 590 Agents national policy simulation, developed Gobang within 800 seconds—existence and bibliographic record verified, 590 number in abstract original	✓	P9
Tier two ★★★★☆ Key Driver					
P4	The Last AI Built by Humans: Toward Genuine Recursive Self-Improvement	arXiv:2609.11873v2 (v1 2026-09-10 / v2 2026-09-15)	RSI five-stage roadmap (execution→strategy→experience acquisition→environment adaptation→recursive meta-improvement) + Headroom-Closed Index (HCI) ruler	✓	P10
P5	Is Recursive Self-Improvement Really Here?	CACM BLOG@CACM , 2026-07-06, journalist Logan Kugler (news column, not peer-reviewed)	Distinguishing "tactical recursion vs strategic recursion," citing MIT CSAIL view: data collection and physical observability are RSI's fundamental bottleneck—the dichotomy and CSAIL view column retellings	✓ (existence; 细分 claims not independently rechecked)	P10 / P14
P6	The Cost of Consensus:	arXiv:2605.00914 (venue correction:	Homogeneous debate costs 2.1–3.4× more tokens, isolated self-	●	P0 / P36

Rank	Title (corrected)	Source and number	Core claim (with evidence qualification)	Verification status	UDOS dialogue paper
	Isolated Self-Correction Prevails Over Unguided Homogeneous Multi-Agent Debate	list's "ACM CAIS 2026-05" not verified)	correction superior; "sycophantic compliance adoption rate 85.5%," "consensus collapse oracle gap 32.3pp" not verified in public abstract , self-reported		
P7	AgentCollabBench: Diagnosing When Good Agents Make Bad Collaborators	arXiv:2605.08647v1 (2026-05-09; "ICML 2026" acceptance not verified)	900 human-verified tasks (software engineering/DevOps/data engineering); "multi-Agent reliability essentially a structural problem" author's conclusive interpretation	●	P9
P8	RSIAgent: Autonomous Exploration for Recursive Self-improvement in New Environments	arXiv:2609.15364 (2026-09, Aether AI)	training-free multi-Agent, curriculum/actor/verifier three-role closed loop, broad-then-deep experience self-exploration; "surpasses GPT-6 Astra" vendor/media claim	✓ (existence; benchmarks not independently rechecked)	P10
P9	SCD v3.1: Structured Contextual Distillation — Deterministic External State Protocol	Zenodo, DOI 10.5281/zenodo.17787619	RFC 8785 JSON canonicalization, SHA-256 integrity chain, turn-based versioning, constitutional governance layer; "Gemini→Claude→Gemini cross-vendor state retention" not verified in public excerpts	●	P14
P10	Adaptive Heterogeneous Multi-Agent Debate for Enhanced Educational and Factual Reasoning	Springer system, Scilit indexed (formal volume/DOI page not located)	A-HMAD (diverse experts + dynamic routing + learned consensus) exists; "heterogeneous over homogeneous improves 4–6% accuracy" self-reported	✓ (existence; 4–6% self-reported)	P9
Tier three ★★★☆☆ Important Supplement					
P11	Recuris: Recursive Experiential–Working Memory Evolution for Long-Horizon Agent Harnesses	arXiv:2608.24876v1 (2026-08-25)	Long-horizon memory recursive evolution; "improved 35 of 37 model-benchmark pairs," tau-bench GPT-5.6 Sol +17.8 / Claude Opus 5 87.9% all self-reported	●	P10 / P14
P12	Phase Transition for Budgeted Multi-Agent Synergy	arXiv:2601.17311v1 (2026-01)	Fixed-point analysis: deep b-ary tree's sharp amplification–collapse phase transition, scalar	✓	P9 / P14

Rank	Title (corrected)	Source and number	Core claim (with evidence qualification)	Verification status	UDOS dialogue paper
			α_p deciding weak signals amplified or washed to random		
P13	When More Agents Hurt: Generalized Amdahl Bounds for Speculative Parallelism in Agentic Software Pipelines	TechRxiv, DOI 10.36227/techrxiv.177220351.10957097/v1 (2026-02; do not confuse with same-named Google DeepMind blog)	Generalized Amdahl lower bound constrained by DAG critical path, beyond critical speculative width selection/merge and serial bottleneck reverse acceleration	✓	P0 / P14
P14	Minority Sentinel: When to Overturn Majority Voting in Multi-Agent LLM Debates	arXiv:2606.29270 (2026-06)	LightGBM meta-classifier identifying moments to overturn majority voting; "stable Flip Precision 81.2%, six-dataset Net Gain positive" self-reported	●	P9 / P14
P15	Verification Capacity Saturation: Three Levers, One Default	AgentPatterns.ai (2026-08-09, Addy Osmani, industry blog/engineering experience, not peer-reviewed)	Verification capacity sets quality ceiling; levers: reduce generation rate / WIP=1 / Little's Law	✓ (existence; empirical claims)	P0 / P14
P16	Portable Agent Memory: A Protocol for Provenance-Verified Memory Transfer Across Heterogeneous LLM Agents	arXiv:2605.11032v1 (2026-05; title correction: Provenance-Verified / Heterogeneous LLM Agents , not "Cryptographically-Verified / Heterogeneous AI Agents")	M=(E,S,P,W,I) five-component memory + Merkle-DAG provenance + capability tokens + injection-resistant rehydration	✓	P14
P17	The Reflexivity Boundary: Why Recursive Self-Improvement Requires External Stabilization	Clue claims Zenodo; two search rounds no such record	"Closed single-observer systems cannot internally produce an L4 jump"— no evidence found , only thematically adjacent self-reference/RSI papers, no same-named entry	■	P36 (must disclose honestly)
P18	Stochasticity in Agentic Evaluations: Quantifying Inconsistency with Intraclass Correlation	arXiv:2512.06710v1 (2025-12-07)	Using intraclass correlation (ICC) decomposing between-query / within-query variance, on GAIA, FRAMES giving ICC intervals	✓	P14
P19	The Computational Boundary of	arXiv:2605.27381v1 (2026-05)	Limited internal self-modification kept within C(A), stabilizing	✓	P10 / P14

Rank	Title (corrected)	Source and number	Core claim (with evidence qualification)	Verification status	UDOS dialogue paper
	Inference: Capability Internalization, Training, and the Turing Jump		revisions governed by jump A' via relativized limit lemma		
P20	Governance by Construction for Generalist Agents	IBM Research / arXiv:2605.10555 (2026-05; venue correction: list's "ACM CAIS 2026-05" not verified)	Five structural checkpoints: Intent Guard / Playbook / Tool Guide / Tool Approvals / Output Formatter	✓	P14

Section summary: among tier one's three papers, P3 is the only ✓ formal peer-reviewed hard foundation; P2's paradigm vision important but 🟡 (single-author thin preprint); P1 directly 🟡. Tier two mostly "tools/ protocols/negative results," venues and numbers each discounted. Tier three's P17 also 🟡—**both empty titles land in "sounds grandest" positions, itself a signal: the more manifesto-like a title, the more it must be checked.**

3 The 10 Industry Milestones Influence Ranking Table

Industry milestones ranked not by "product usability," but by "whether it changed multi-Agent systems' default substrate or RSI's public narrative"; evidence axis likewise overlaid.

Table 2 10 industry milestones • influence tier × evidence three states

Rank	Milestone (corrected)	Source and number	Core claim (with evidence qualification)	Verification status	UDOS dialogue paper
I1	MCP 2026-07-28 largest revision + A2A v1.0 + same governance framework	MCP official blog "The 2026-07-28 Specification"; A2A official "A2A Protocol Ships v1.0"; Linux Foundation AAIF	stateless core, round-robin load balancing over ordinary HTTP, protocol-level sessions removed, A2A Signed Agent Card (cryptographic signatures) true ; "150+ supporting organizations" not directly located (AAIF 2026-05 claims total membership about 190)	🟡 (standard release/ vendor disclosure)	P0 / P1
I2	OpenAI GPT-5.3-Codex "instrumental in creating itself"	OpenAI official "Introducing GPT-5.3-Codex" (2026-02-05)	Site verbatim "GPT-5.3-Codex is our first model that was instrumental in creating itself," House hearing testimony also retells— company self-statement, not independent reproduction	🟡 (company disclosure)	P10 / P14
I3	Volcengine four-layer Agent product matrix	volcengine.com official product list	AgentKit, HiAgent 3.0, ArkClaw enterprise workbench, TRAE CN, Coze zero-code platform—five products genuinely exist; "four-layer matrix" editorial induction	✅	P1
I4	Warp-Cortex single-card 100 concurrent Agents (2.2GB VRAM)	arXiv:2601.01298 (same paper P2)	100 concurrent @2.2GB single-author self-reported, no third-party reproduction	🟡 (paper self-reported)	P1 / P14
I5	Anthropic R&D Automation Index	Anthropic official "Measurements for understanding the pace of AI development inside frontier labs" (2026-09-17/18)	About 26% own AI R&D reaches AL4 "leading," >90% at least AL3, about 30,000 internal Agents concurrent, at 2 months leading share <1%— Anthropic self-assessment, no external independent verification	🟡 (company disclosure)	P10
I6	Handoff / Agent-as-Tool orchestration primitive standardization	OpenAI Agents SDK official docs; Microsoft AutoGen official docs	Handoff transfers conversation ownership, receiver faces user directly; Agent-as-Tool orchestrator retains control, sub-Agent returns strings—both official docs clearly distinguish	✅	P9 / P14

Rank	Milestone (corrected)	Source and number	Core claim (with evidence qualification)	Verification status	UDOS dialogue paper
I7	Chinese Agent platform ecosystem mature (IDC caliber: Volcengine leading on both sides)	IDC "China Agent Development Platform Market Share, 2025" (2026-06-12), retold via Sina Finance/36Kr	2025 private market about 1.75 billion yuan, Volcengine 17.8% private first, Coze 19.3% public cloud first— media retelling, IDC original PDF not consulted	🟡 (third-party report/media retelling)	P1
I8	DeepSeek Harness / Cordis reversible plugin system	deepseek.com/harness, deepseekharness.dev; v0.1 developer preview 2026-08-13, MIT license	"Everything is a plugin," Cordis plugin kernel, MIT license, derived from Koishi ecosystem (about four years/4000+ plugins lineage claims, not itemized), dependency injection + revertible effects	✅	P0 / P14
I9	Epoch AI's AL0–AL5 automation grading adopted by industry	Anthropic official; media like Jiqizhixin	Anthropic official clearly adopts Epoch AI automation grading AL0→AL5, AL4=AI leading; Epoch AI an independent nonprofit	✅	P10
I10	Agentic AI Foundation (AAIF) established, MCP and A2A unified governance	Linux Foundation / AAIF official announcement	AAIF established by Linux Foundation 2025-12-09 (founding projects include MCP, goose, AGENTS.md); A2A only merged 2026-08-17 —direction true, timeline must be corrected	✅	P0 / P1

Section summary: industry's hardest substrate is **protocols and orchestration primitives** (I1's stateless MCP, I6's Handoff standardization, I8's reversible plugin kernel), softest narrative is **company self-reported RSI numbers** (I2, I5). Market share type (I7) uniformly read as retelling. Note I1 and I10 are two faces of the same governance line: protocol substrate ✅🟡 true, but "150+ organizations" and "same governance" exact calibers discounted.

4 MegaAgent Conflict Resolution: Same Paper, Not Two

The v7.7.1 list and v7.6 ledger once seemed to describe two milestones—one called ...A Practical Framework..., one ...Without Predefined SOPs.... **Resolution: same paper, two versions of early preprint title and formal publication title.**

- **Number unique:** indeed **arXiv:2408.09955** (v1 submitted 2024-08-19, v2 2024-08-20, v3 2025-05-29). The list's number correct.
- **Title evolution:**

- Early v1/v2 (2024-08): MegaAgent: A Practical Framework for Autonomous Cooperation in Large-Scale LLM Agent Systems.
- v3 (2025-05-29) and ACL formal publication: MegaAgent: A Large-Scale Autonomous LLM-based Multi-Agent System Without Predefined SOPs.
- **Formal source:** Findings of ACL 2025, Anthology ID **2025.findings-acl.259**, pp. **4998–5036**, Vienna, 2025-07, DOI 10.18653/v1/2025.findings-acl.259.
- **"590 Agent" attribution:** belongs to **this paper**. arXiv abstract and ACL abstract original both state "scaling up to **590 agents in a national policy simulation**." Whether paired with old or new title not wrong, but **uniformly cite formal title**.
- **UDOS citation caliber:** uniformly write ...Without Predefined SOPs + ACL 2025 Findings source, parenthetically noting early alias ...A Practical Framework....

Section summary: do not count MegaAgent as two literature items occupying two slots. It is tier one's only  formal review hard foundation—the 590-Agent SOP-free autonomy fact corresponds only to this paper.

5 No-Evidence Section: Two Empty Titles + a Batch "Exists But Caliber Must Be Corrected"

5.1 Genuinely No Evidence (■, forbidden to cite as fact)

- **P1 Toward Reliable Collective Intelligence in LLM-MAS:** list claims IEEE Xplore 2026-05. Two rounds of different keyword searches (including exact phrase and IEEE restriction) both missed—returned multiple thematically close LLM-MAS surveys (e.g. arXiv:2502.01714 etc.), **none same-named, nor IEEE Xplore 2026-05 record**. In writing only mark [CITATION NEEDED] .
- **P17 The Reflexivity Boundary: Why Recursive Self-Improvement Requires External Stabilization:** list claims Zenodo. Two search rounds (Zenodo + L4 + single observer + external stabilization) both missed—returned thematically adjacent self-reference/RSI papers (e.g. arXiv: 2607.04277 Self-Reference in LLMs), **none same-named**. Likewise only mark [CITATION NEEDED] .

5.2 Exists But Caliber Discrepant (not ■, but must write per corrected caliber)

Item	List claim	Corrected actual
P6 venue / numbers	ACM CAIS 2026-05; 85.5%, 32.3pp	Actually arXiv:2605.00914 preprint ; 85.5%/32.3pp not verified in public abstract, mark self-reported
P7 venue	ICML 2026	Actually arXiv:2605.08647 , "ICML 2026" no acceptance evidence found
P2 number	98% context compression no semantic loss	Abstract only claims witness-complex sparsification preserves context manifold persistent homology features, 98% not appearing
P20 venue	ACM CAIS 2026-05	Actually IBM Research / arXiv:2605.10555
I1 number	150+ supporting organizations	Not directly located; AAIF 2026-05 claims total membership about 190
I7 number	17.8% / 19.3%	IDC report media retelling , original PDF not consulted

Section summary: ■ only P1, P17 two papers, their "no evidence" itself part of the ranking conclusion—cannot place in foundational tier as fact because it sounds important. Remaining 🟡 items not false, they are "numbers not yet at directly citable strength," write per Table 1/Table 2 qualifiers.

6 Dross, Failure Boundaries and Falsifiable Conditions

This ranking not only tells readers "who is important," but also centrally marks **which popular narratives have been falsified by this batch of literature itself or must be qualified**.

- **"Multi-Agent debate necessarily more accurate"—falsified by P6.** P6's title itself Isolated Self-Correction Prevails Over Unguided Homogeneous Multi-Agent Debate: unguided homogeneous debate produces sycophantic compliance, consensus collapse, costs 2.1–3.4× more tokens instead inferior to isolated self-correction. Conclusion: debate should be "heterogeneous + guided" (see P10), not "more people stronger."
- **"Scaling Laws seamlessly extrapolate to multi-Agent"—falsified by P12, P13.** P12 proves budgeted multi-Agent synergy has a sharp amplification–collapse **phase transition** (past the critical point weak signals washed to random); P13 gives generalized **Amdahl lower bound** (DAG critical path serial bottleneck reverses acceleration past critical width). Conclusion: Agent count not more-is-better, no "mindless headcount" scaling curve.
- **"RSI already happening"—must distinguish P5's tactical vs strategic recursion.** P5 (news column caliber) suggests: what we currently see (I2 GPT-5.3-Codex self-statement, I5 Anthropic 26% AL4) mostly **tactical-level** AI participation in building/evaluating the next generation; P4's five-stage roadmap and P19's Turing Jump theory show genuine **strategic recursion (recursive meta-improvement)** not yet established. In citing RSI signals one must not write tactical recursion as strategic recursion achieved.
- **■ items must not be cited as fact.** P1, P17 even if placed in some tier, only mentionable in "to-be-verified literature" sense, their claims (LLM-MAS coordination survey, closed single observer cannot L4) uniformly mark [CITATION NEEDED].
- **All 🟡 numbers attach falsifiable conditions:** 85.5%/32.3pp (P6), 98% (P2), 35/37 (P11), 81.2% (P14), 4–6% (P10), 26%/>90%/about 30,000 (I5), 17.8%/19.3% (I7)—if third-party reproduction or vendor originals after publication significantly differ from self-reported values, should change to ■ and correct this paper.

Section summary: the frontier's most valuable thing is not "how many more points gained," but this batch's **negative results and boundaries**—debate has ceilings, headcount has phase transitions, RSI needs layering. Citing these as conclusions is more stable than citing any single benchmark.

7 References / Evidence Links

The following links all genuinely located during EVIDENCE_LEDGER_v771.md verification; unverified ones uniformly not listed, not fabricated. DOI / arXiv numbers / pages per this table.

Papers (P1–P20)

1. P2 Warp-Cortex: <https://arxiv.org/abs/2601.01298>

2. P3 MegaAgent (formal version): <https://aclanthology.org/2025.findings-acl.259/> ; early alias <https://arxiv.org/abs/2408.09955>
3. P4 The Last AI Built by Humans: <https://arxiv.org/abs/2609.11873>
4. P5 Is Recursive Self-Improvement Really Here? (BLOG@CACM): <https://cacm.acm.org/>
5. P6 The Cost of Consensus: <https://arxiv.org/abs/2605.00914>
6. P7 AgentCollabBench: <https://arxiv.org/abs/2605.08647>
7. P8 RSIAgent: <https://arxiv.org/pdf/2609.15364>
8. P9 SCD v3.1 (Zenodo): <https://zenodo.org/records/17787619>
9. P10 Adaptive Heterogeneous Multi-Agent Debate (Scilit): <https://www.scilit.com/publications/cb34d8466c8b48837a98277cdde57aea>
10. P11 Recuris: <https://arxiv.org/abs/2608.24876>
11. P12 Phase Transition for Budgeted Multi-Agent Synergy: <https://arxiv.org/html/2601.17311v1>
12. P13 Generalized Amdahl Bounds (TechRxiv / ESS Open Archive): <https://essopenarchive.org/doi/full/10.36227/techrxiv.177220351.10957097/v1>
13. P14 Minority Sentinel: <https://arxiv.org/html/2606.29270>
14. P15 Verification Capacity Saturation (AgentPatterns.ai): <https://agentpatterns.ai/verification/verification-capacity-saturation/>
15. P16 Portable Agent Memory: <https://arxiv.org/html/2605.11032v1>
16. P18 Stochasticity in Agentic Evaluations: <https://arxiv.org/abs/2512.06710>
17. P19 The Computational Boundary of Inference / Turing Jump: <https://arxiv.org/html/2605.27381v1>
18. P20 Governance by Construction (IBM Research): <https://research.ibm.com/publications/governance-by-construction-for-generalist-agents>
19. P1, P17: **no evidence, no evidence links** (see Section 5.1).

Industry milestones (I1–I10)

1. I1 MCP 2026-07-28: <https://blog.modelcontextprotocol.io/posts/2026-07-28/> ; A2A v1.0: <https://a2a-protocol.org/latest/announcing-1.0/>
2. I2 GPT-5.3-Codex: <https://openai.com/index/introducing-gpt-5-3-codex/>
3. I3 Volcengine product matrix: <https://www.volcengine.com/> ; <https://www.volcengine.com/product/list>
4. I4 Warp-Cortex single-card concurrency: <https://arxiv.org/abs/2601.01298>
5. I5 Anthropic R&D Automation Index: <https://www.anthropic.com/institute/measuring-pace-of-ai-development> ; <https://www.anthropic.com/institute/recursive-self-improvement>
6. I6 Handoff / Agent-as-Tool: <https://openai.github.io/openai-agents-python/handoffs/> ; <https://microsoft.github.io/autogen/stable/user-guide/core-user-guide/design-patterns/handoffs.html>
7. I7 IDC caliber (media retelling): <https://finance.sina.com.cn/roll/2026-06-17/doc-inictaet3901961.shtml> ; <https://finance.sina.com.cn/cj/2026-06-25/doc-inierayt7754685.shtml>
8. I8 DeepSeek Harness / Cordis: <https://www.deepseek.com/harness/en/> ; <https://deepseekharness.dev/> ; <https://cloud.tencent.com/developer/article/2727482>
9. I9 Epoch AI AL0–AL5 adopted: <https://www.anthropic.com/institute/measuring-pace-of-ai-development>

10. I10 AAIIF establishment and A2A merger: <https://aaif.io/press/linux-foundation-announces-the-formation-of-the-agentic-ai-foundation-aaif-anchored-by-new-project-contributions-including-model-context-protocol-mcp-goose-and-agents-md/> ; <https://aaif.io/blog/a2a-joins-aaif>

Evidence Discipline and Reproduction Notes

- This paper is an **editorial influence ranking + evidence verification** paper, no new experiments; all external facts take `EVIDENCE_LEDGER_v771.md` as the sole source, verification base date 2026-09-21.
- Influence tiers (★★★★★/★★★★☆/★★★☆☆) are **editorial judgment**, not citation metrics; three-state (✓/●/■) is **verifiable fact**, the two not replacing each other.
- ● numbers uniformly used per "paper self-reported / company disclosed / media retelling"; ■ (P1, P17) uniformly [CITATION NEEDED], not written as fact.
- MegaAgent uniformly cites formal title ...Without Predefined SOPs + ACL 2025 Findings (2025.findings-acl.259, pp.4998–5036), not split into two.
- Before submission all ● items should have full text/official originals opened item by item for secondary recheck; if reproduced values significantly differ from self-reported, change per Section 6 falsifiable conditions.